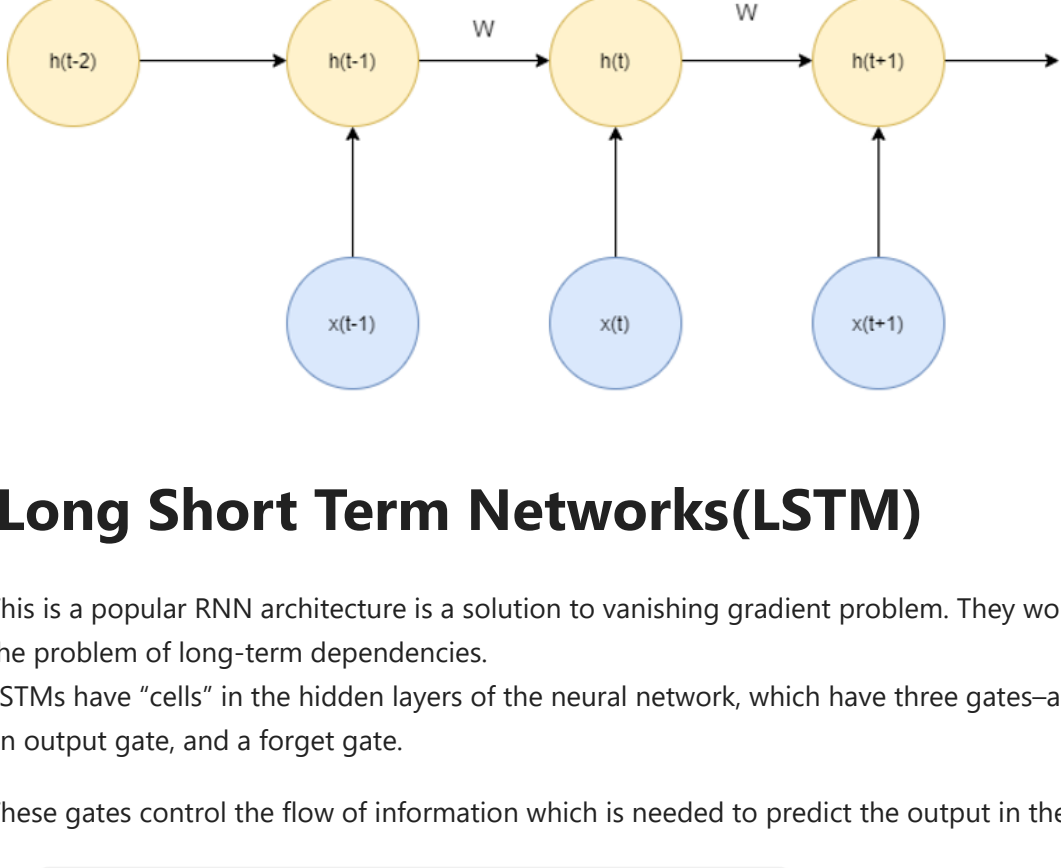


Recurrent Neural Networks & Day Trading (LSTM)

1. Recurrent Neural Network

- recurrent neural network store the output activations from one or more of the layers of the network, often called **hidden layers**
 - The next time we feed an input to the network, we include the previously-stored outputs as additional inputs.
- RNNs are quite powerful but, they suffer from **Vanishing gradient**
- It hinders them from long term information, hence limiting RNN with poor results.
- Instead, we use a better variation of RNNs: **Long Short Term Networks(LSTM)**.



2. Long Short Term Networks(LSTM)

- This is a popular RNN architecture is a solution to vanishing gradient problem. They work to address the problem of long-term dependencies.
 - LSTMs have “cells” in the hidden layers of the neural network, which have three gates—an input gate, an output gate, and a forget gate.
- These gates control the flow of information which is needed to predict the output in the network.



3. Code Execution

```
In [ ] :  
# Importing the libraries  
import numpy as np  
import matplotlib.pyplot as plt  
plt.style.use('fivethirtyeight')  
import pandas as pd  
from sklearn.preprocessing import MinMaxScaler  
from keras.models import Sequential  
from keras.layers import Dense, LSTM, Dropout, GRU, Bidirectional  
from keras.optimizers import SGD  
import math  
from sklearn.metrics import mean_squared_error
```

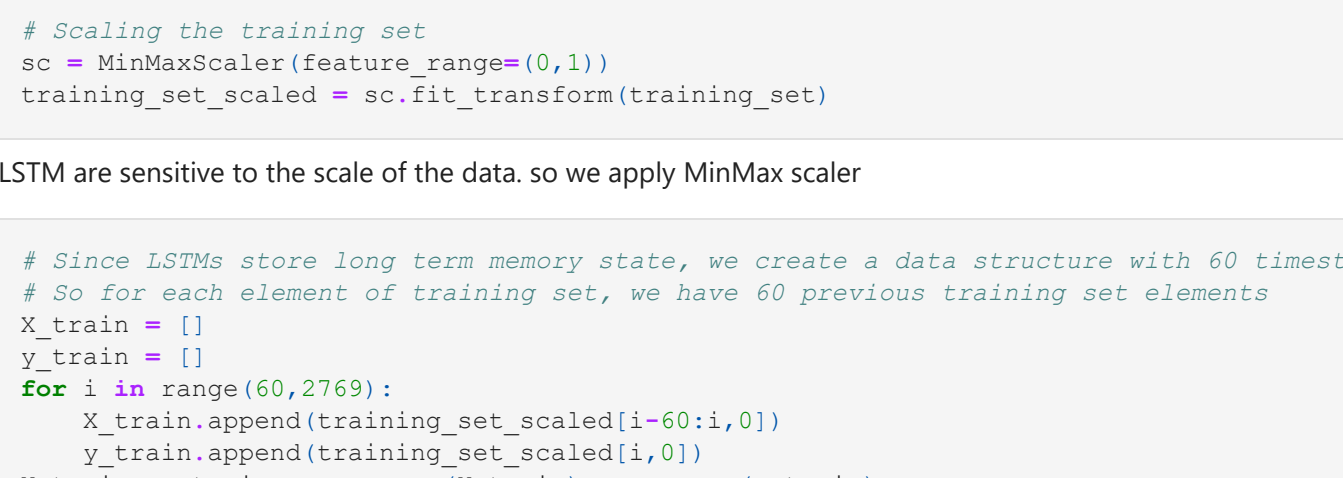
```
In [ ] :  
# Some functions to help out with  
def plot_predictions(test,predicted):  
    plt.plot(test, color='red',label='Real IBM Stock Price')  
    plt.plot(predicted, color='blue',label='Predicted IBM Stock Price')  
    plt.title('IBM Stock Price Prediction')  
    plt.xlabel('Time')  
    plt.ylabel('IBM Stock Price')  
    plt.legend()  
    plt.show()  
  
def return_rmse(test,predicted):  
    rmse = math.sqrt(mean_squared_error(test, predicted))  
    print("The root mean squared error is {}".format(rmse))
```

```
In [ ] :  
# First, we get the data  
dataset = pd.read_csv('https://github.com/koshtinikhilesh/NeuralNetwork_Project/blob/master/IBM_Stock_Price.csv')  
dataset.head()
```

	Open	High	Low	Close	Volume	Name
Date						
2006-01-03	82.45	82.55	80.81	82.06	11715200	IBM
2006-01-04	82.20	82.50	81.33	81.95	9840600	IBM
2006-01-05	81.40	82.90	81.00	82.50	7213500	IBM
2006-01-06	83.95	85.03	83.41	84.95	8197400	IBM
2006-01-09	84.10	84.25	83.38	83.73	6858200	IBM

```
In [ ] :  
# Checking for missing values  
training_set = dataset['2016'].iloc[:,1:2].values  
test_set = dataset['2017'].iloc[:,1:2].values
```

```
In [ ] :  
# We have chosen 'High' attribute for prices. Let's see what it looks like  
dataset["High"][:'2016'].plot(figsize=(16,4),legend=True)  
dataset["High"][:'2017'].plot(figsize=(16,4),legend=True)  
plt.legend(['Training set (Before 2017)', 'Test set (2017 and beyond)'])  
plt.title('IBM stock price')  
plt.show()
```



```
In [ ] :  
# Scaling the training set  
sc = MinMaxScaler(feature_range=(0,1))  
training_set_scaled = sc.fit_transform(training_set)
```

LSTM are sensitive to the scale of the data. so we apply MinMax scaler

```
In [ ] :  
# Since LSTMs store long term memory state, we create a data structure with 60 timesteps  
# So for each element of training set, we have 60 previous training set elements  
X_train = []  
y_train = []  
for i in range(60,2769):  
    X_train.append(training_set_scaled[i-60:i,0])  
    y_train.append(training_set_scaled[i,0])  
X_train, y_train = np.array(X_train), np.array(y_train)
```

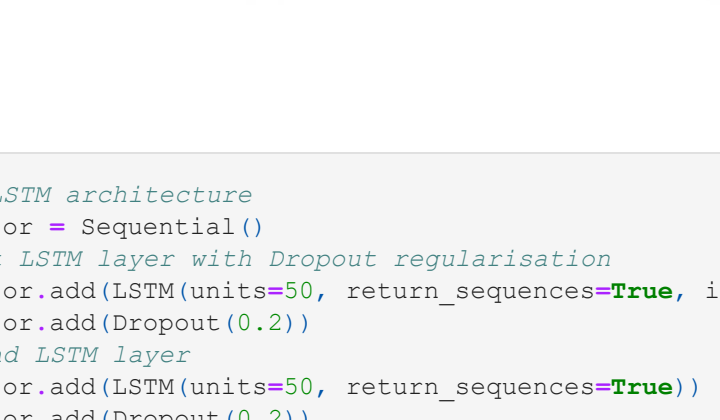
- We have taken 60 as our time step value.
- n no of time_steps in LSTM = n no of previous training elements

```
In [ ] :  
# Reshaping X_train for efficient modelling  
X_train = np.reshape(X_train, (X_train.shape[0],X_train.shape[1],1))
```

- We use reshape() funtion to change the shape of the array. Here we added 1 to convert 2D array to 3D array.
- 2D array provide medium complexity but 3D array provide maximum complexity.

3. Stacked LSTM Model

The Stacked LSTM is nothing but an LSTM Model with multiple LSTM layers. The LSTM layer gives a sequential output to the next LSTM layer.



```
In [ ] :  
# The LSTM architecture  
regressor = Sequential()  
# First LSTM layer with Dropout regularisation  
regressor.add(LSTM(units=50, return_sequences=True, input_shape=(X_train.shape[1],1)))  
regressor.add(Dropout(0.2))  
# Second LSTM layer  
regressor.add(LSTM(units=50, return_sequences=True))  
regressor.add(Dropout(0.2))  
# Third LSTM layer  
regressor.add(LSTM(units=50, return_sequences=True))  
regressor.add(Dropout(0.2))  
# Fourth LSTM layer  
regressor.add(LSTM(units=50))  
regressor.add(Dropout(0.2))  
# The output layer  
regressor.add(Dense(units=1))
```

```
# Compiling the RNN  
regressor.compile(optimizer='rmsprop',loss='mean_squared_error')  
# Fitting to the training set  
regressor.fit(X_train,y_train,epochs=50,batch_size=32)
```

```
Epoch 1/50  
85/85 [=====] - 22s 134ms/step - loss: 0.0236  
Epoch 2/50  
85/85 [=====] - 11s 129ms/step - loss: 0.0108  
Epoch 3/50  
85/85 [=====] - 11s 130ms/step - loss: 0.0081  
Epoch 4/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0074  
Epoch 5/50  
85/85 [=====] - 13s 143ms/step - loss: 0.0064  
Epoch 6/50  
85/85 [=====] - 12s 145ms/step - loss: 0.0057  
Epoch 7/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0053  
Epoch 8/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0046  
Epoch 9/50  
85/85 [=====] - 12s 144ms/step - loss: 0.0043  
Epoch 10/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0038  
Epoch 11/50  
85/85 [=====] - 11s 134ms/step - loss: 0.0037  
Epoch 12/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0034  
Epoch 13/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0033  
Epoch 14/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0030  
Epoch 15/50  
85/85 [=====] - 11s 133ms/step - loss: 0.0030  
Epoch 16/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0029  
Epoch 17/50  
85/85 [=====] - 14s 159ms/step - loss: 0.0028  
Epoch 18/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0028  
Epoch 19/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0026  
Epoch 20/50  
85/85 [=====] - 11s 133ms/step - loss: 0.0025  
Epoch 21/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0025  
Epoch 22/50  
85/85 [=====] - 13s 148ms/step - loss: 0.0023  
Epoch 23/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0024  
Epoch 24/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0023  
Epoch 25/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0022  
Epoch 26/50  
85/85 [=====] - 11s 129ms/step - loss: 0.0022  
Epoch 27/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0022  
Epoch 28/50  
85/85 [=====] - 11s 130ms/step - loss: 0.0020  
Epoch 29/50  
85/85 [=====] - 11s 130ms/step - loss: 0.0020  
Epoch 30/50  
85/85 [=====] - 11s 130ms/step - loss: 0.0021  
Epoch 31/50  
85/85 [=====] - 11s 130ms/step - loss: 0.0019  
Epoch 32/50  
85/85 [=====] - 11s 130ms/step - loss: 0.0018  
Epoch 33/50  
85/85 [=====] - 11s 133ms/step - loss: 0.0017  
Epoch 34/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0018  
Epoch 35/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0017  
Epoch 36/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0017  
Epoch 37/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0017  
Epoch 38/50  
85/85 [=====] - 12s 146ms/step - loss: 0.0018  
Epoch 39/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0017  
Epoch 40/50  
85/85 [=====] - 11s 134ms/step - loss: 0.0017  
Epoch 41/50  
85/85 [=====] - 11s 135ms/step - loss: 0.0016  
Epoch 42/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0016  
Epoch 43/50  
85/85 [=====] - 11s 133ms/step - loss: 0.0015  
Epoch 44/50  
85/85 [=====] - 11s 132ms/step - loss: 0.0015  
Epoch 45/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0015  
Epoch 46/50  
85/85 [=====] - 11s 134ms/step - loss: 0.0015  
Epoch 47/50  
85/85 [=====] - 12s 136ms/step - loss: 0.0016  
Epoch 48/50  
85/85 [=====] - 12s 136ms/step - loss: 0.0014  
Epoch 49/50  
85/85 [=====] - 11s 133ms/step - loss: 0.0015  
Epoch 50/50  
85/85 [=====] - 11s 131ms/step - loss: 0.0014
```


Used 4 LSTM layers for the model and the optimizer is Rmsprop

```
In [ ] :  
# Now to get the test set ready in a similar way as the training set.  
# The following has been done so farst 60 entires of test set have 60 previous values  
# 'High' attribute data for processing  
dataset_total = pd.concat((dataset["High"][:'2016'],dataset["High"][:'2017']),axis=0)  
inputs = dataset_total[len(dataset_total)-len(test_set) : 60:].values  
inputs = inputs.reshape(-1,1)  
inputs = sc.transform(inputs)
```

```
In [ ] :  
# Preparing X_test and predicting the prices  
X_test = []  
for i in range(60,311):  
    X_test.append(inputs[i-60:i,0])  
X_test = np.array(X_test)  
X_test = np.reshape(X_test, (X_test.shape[0],X_test.shape[1],1))  
predicted_stock_price = regressor.predict(X_test)  
predicted_stock_price = sc.inverse_transform(predicted_stock_price)
```

```
8/8 [=====] - 2s 34ms/step
```

```
In [ ] :  
# Visualizing the results for LSTM  
plot_predictions(test_set,predicted_stock_price)
```



```
In [ ] :  
# Evaluating our model  
return_rmse(test_set,predicted_stock_price)
```

The root mean squared error is 2.5942884190823072.

4. Conclusion

1. Implemented LSTM using Keras Library.
2. The data implemented in this model is time series data. Therefore, the data is divided with respect to year value.
3. Model predicts the high value of the dataset.
4. Since LSTM is sensitive to scale, we have implemented MinMaxScaler in order to convert all high values in our dataset to range of 0 to 1.
5. Our Training Data is the entire array of data before year = 2016.
6. Our Testing Data contains all the values after year = 2017.
7. In our LSTM Model we are using 60 time_steps.
8. We use reshape() funtion to change the shape of the array. Here we added 1 to convert 2D array to 3D array.
9. 2D array provide medium complexity but 3D array provide maximum complexity.
10. Used 4 LSTM layers for the model and the optimizer is Rmsprop instead of Adam because of better RMSE value close to 0 -> 0.5
11. RMSPROP_RMSE value = 2.5 ADAM_RMSE value = 5.2